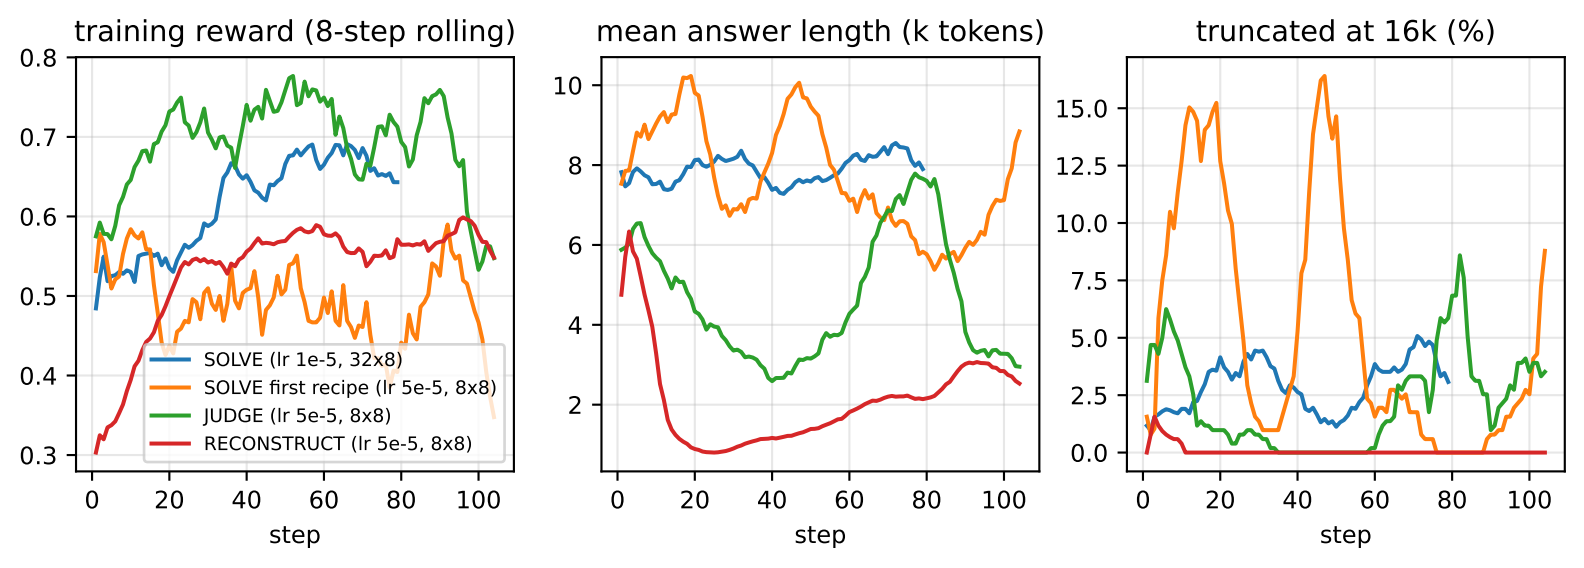
training reward (8-step rolling)
mean answer length (k tokens)
truncated at 16k (%)
step
step
step
SOLVE (lr 1e-5, 32x8)
SOLVE first recipe (lr 5e-5, 8x8)
JUDGE (lr 5e-5, 8x8)
RECONSTRUCT (lr 5e-5, 8x8)